

Robustifying Profile Information Propagation in Profile-Guided Optimization



SAPIENZA
UNIVERSITÀ DI ROMA

Table of Contents

1. Profile-Guided Optimization
2. Practical Obstacles
3. State of the Art
4. Profile Mishandling Research Gap
5. Research Proposal
6. Evaluation of the Proposed Directions
7. Impacts and Benefits

Performance is Critical

Performance is Critical

- Achieving optimal performance is of critical importance today.
 - A 1% performance improvement can save millions of dollars annually in large-scale deployments.

Performance is Critical


- Achieving optimal performance is of critical importance today.
 - A 1% performance improvement can save millions of dollars annually in large-scale deployments.
- **Profile-Guided Optimization** is a way to achieve great performance
 - Optimization process **tailored** to the workload of the compiled program
 - **Profile** = Execution frequencies of program regions
 - Eliminates the need to infer such information
 - Perfect for high-load applications

Profile-Guided Optimization


Profile-Guided Optimization

Source program


Profile-Guided Optimization



Profile-Guided Optimization




Profile-Guided Optimization




Profile-Guided Optimization



Profile-Guided Optimization



Profile-Guided Optimization




Profile Life-cycle

Profile Life-cycle

- Profiles traverse three main phases within the PGO workflow


Profile Life-cycle

- Profiles traverse three main phases within the PGO workflow
 - Collection:** The profile is collected and persisted as a file




Profile Life-cycle

- Profiles traverse three main phases within the PGO workflow
 - **Collection:** The profile is collected and persisted as a file
 - **Mapping:** The profile is represented as **metadata** attached to program constructs



Profile Life-cycle

- Profiles traverse three main phases within the PGO workflow
 - **Collection:** The profile is collected and persisted as a file
 - **Mapping:** The profile is represented as **metadata** attached to program constructs
 - **Usage:** The metadata is used to guide optimizing transformations



Practical Obstacles

- For an ideal PGO application



Profile must be **complete** and **accurate** in each phase of its life-cycle

Practical Obstacles

- For an ideal PGO application



Profile must be **complete** and **accurate** in each phase of its life-cycle

- Inaccurate profiles may lead the compiler to make suboptimal optimization decisions.

Practical Obstacles

- For an ideal PGO application



Profile must be **complete** and **accurate** in each phase of its life-cycle

- Inaccurate profiles may lead the compiler to make suboptimal optimization decisions.
- In practice, each phase hides sources of inaccuracies
 - **Sampling** strategies are inaccurate by nature
 - The dynamic nature of software leads to **stale profiles**
 - Optimizing passes contain errors in **metadata propagation** logic

State of the Art

1. Wenlei He, Julián Mestre, Sergey Pupyrev, Lei Wang, and Hongtao Yu. "Profile inference revisited". ↵
2. Amir Ayupov, Maksim Panchenko, and Sergey Pupyrev. "Stale Profile Matching". ↵
3. Youfeng Wu. "Accuracy of Profile Maintenance in Optimizing Compilers". ↵
4. Profile Information Propagation Unittesting: <https://discourse.llvm.org/t/rfc-profile-information-propagation-unittesting/73595> ↵

State of the Art

- Lots of effort to solve problems in the collection and mapping phases

1. Wenlei He, Julián Mestre, Sergey Pupyrev, Lei Wang, and Hongtao Yu. "Profile inference revisited". ↵
2. Amir Ayupov, Maksim Panchenko, and Sergey Pupyrev. "Stale Profile Matching". ↵
3. Youfeng Wu. "Accuracy of Profile Maintenance in Optimizing Compilers". ↵
4. Profile Information Propagation Unittesting: <https://discourse.llvm.org/t/rfc-profile-information-propagation-unittesting/73595> ↵

State of the Art

- Lots of effort to solve problems in the collection and mapping phases
 - [1] Proposes a rectification algorithm to rectify sampled profiles

1. Wenlei He, Julián Mestre, Sergey Pupyrev, Lei Wang, and Hongtao Yu. "Profile inference revisited". ↵
2. Amir Ayupov, Maksim Panchenko, and Sergey Pupyrev. "Stale Profile Matching". ↵
3. Youfeng Wu. "Accuracy of Profile Maintenance in Optimizing Compilers". ↵
4. Profile Information Propagation Unittesting: <https://discourse.llvm.org/t/rfc-profile-information-propagation-unittesting/73595> ↵

State of the Art

- Lots of effort to solve problems in the collection and mapping phases
 - [\[1\]](#) Proposes a rectification algorithm to rectify sampled profiles
 - [\[2\]](#) Proposes an algorithm to adapt stale profiles to newer program versions

1. Wenlei He, Julián Mestre, Sergey Pupyrev, Lei Wang, and Hongtao Yu. "Profile inference revisited". ↵
2. Amir Ayupov, Maksim Panchenko, and Sergey Pupyrev. "Stale Profile Matching". ↵
3. Youfeng Wu. "Accuracy of Profile Maintenance in Optimizing Compilers". ↵
4. Profile Information Propagation Unittesting: <https://discourse.llvm.org/t/rfc-profile-information-propagation-unittesting/73595> ↵

State of the Art

- Lots of effort to solve problems in the collection and mapping phases
 - [\[1\]](#) Proposes a rectification algorithm to rectify sampled profiles
 - [\[2\]](#) Proposes an algorithm to adapt stale profiles to newer program versions
 - [\[3\]](#), [\[4\]](#) only partially address failures in metadata propagation

1. Wenlei He, Julián Mestre, Sergey Pupyrev, Lei Wang, and Hongtao Yu. “Profile inference revisited”. ↵


2. Amir Ayupov, Maksim Panchenko, and Sergey Pupyrev. “Stale Profile Matching”. ↵

3. Youfeng Wu. “Accuracy of Profile Maintenance in Optimizing Compilers”. ↵

4. Profile Information Propagation Unittesting: <https://discourse.llvm.org/t/rfc-profile-information-propagation-unittesting/73595> ↵


Example of Profile Mishandling

```
// Before pass  
if (x > 0) { // Then branch taken 80 times  
    handle_positive(x) 🔥  
} else { // Else branch taken 20 times  
    handle_negative(x) ❄️  
}
```




Example of Profile Mishandling

```
// Before pass  
if (x > 0) { // Then branch taken 80 times  
    handle_positive(x) 🔥  
} else { // Else branch taken 20 times  
    handle_negative(x) ❄️  
}
```



```
// After pass  
if (x ≤ 0) { // Then branch taken 80 times  
    handle_negative(x) 🔥  
} else { // Else branch taken 20 times  
    handle_positive(x) ❄️  
}
```



Profile Mishandling Research Gap

Profile Mishandling Research Gap

- Recently discovered **regressions** could be attributed to profile mishandling

Profile Mishandling Research Gap

- Recently discovered **regressions** could be attributed to profile mishandling
- Profile mishandling **nullifies** solutions aiming at correcting profiles in earlier stages

Profile Mishandling Research Gap

- Recently discovered **regressions** could be attributed to profile mishandling
- Profile mishandling **nullifies** solutions aiming at correcting profiles in earlier stages
- **Non-trivial** solution
 - What does it mean to correctly propagate profiles?
 - No immediate correlation between before and after code
 - **Interaction** between passes needs to be taken into account
 - **Static checks** on control-flow are not enough!
 - **No dedicated tools** to validate propagation logic

Profile Mishandling Research Gap

- Recently discovered **regressions** could be attributed to profile mishandling
- Profile mishandling **nullifies** solutions aiming at correcting profiles in earlier stages
- **Non-trivial** solution
 - What does it mean to correctly propagate profiles?
 - No immediate correlation between before and after code
 - **Interaction** between passes needs to be taken into account
 - **Static checks** on control-flow are not enough!
 - **No dedicated tools** to validate propagation logic



Profile information is **transformed** together with the program, but unlike the program itself, its correctness cannot be directly observed.

Research Proposal



Can profile propagation accuracy be assessed systematically?

Research Proposal



Can profile propagation accuracy be assessed systematically?

- A Methodological and practical framework

Research Proposal



Can profile propagation accuracy be assessed systematically?

- A Methodological and practical framework
- **Black-box** fuzzing strategy
 - Stress-test profile propagation logic in order to expose bugs.
 - Off-the-shelf random program generators

Research Proposal




Can profile propagation accuracy be assessed systematically?


- A Methodological and practical framework
- **Black-box** fuzzing strategy
 - Stress-test profile propagation logic in order to expose bugs.
 - Off-the-shelf random program generators
- **Grey-box** fuzzing strategy
 - Lightweight **feedback mechanism** to guide coverage
 - Based on code and **profile** mutations
 - Uses programs from existing **test-suites** as foundation

Black-Box Fuzzing: Methodology


Black-Box Fuzzing: Methodology




Black-Box Fuzzing: Methodology



Black-Box Fuzzing: Methodology




Black-Box Fuzzing: Methodology



Black-Box Fuzzing: Future Directions

- Improving **automated** bug triaging
 - Classify the large number of found mismatches
 - Avoid reporting multiple instances of the same issue
 - Hard to do due to complexity of randomly generated programs
- Evaluate the framework on **real-world** applications
 - Complete the validation of the methodology
 - SPEC CPU 2017 to evaluate the performance impact of bug fixes

Grey-Box Fuzzing: Methodology



- Mutate both IR and profile metadata
- Instrument optimization passes to obtain profile-aware coverage
 - IR transformations
 - Profile-update actions
- Coverage guides exploration toward new optimization behaviors
- Profile mutations exercise rarely tested PGO scenarios

Evaluation of the Proposed Directions

Evaluation of the Proposed Directions

- **LLVM** compiler infrastructure as evaluation target

Evaluation of the Proposed Directions

- **LLVM** compiler infrastructure as evaluation target
- Can the proposed methodologies:

Evaluation of the Proposed Directions


- **LLVM** compiler infrastructure as evaluation target
- Can the proposed methodologies:
 - Find new **bugs**? If so, how many?



Bugs

Evaluation of the Proposed Directions

- **LLVM** compiler infrastructure as evaluation target
- Can the proposed methodologies:
 - Find new **bugs**? If so, how many?
 - **Improve the performance** of the generated binaries?



Bugs

Performance

Evaluation of the Proposed Directions

- **LLVM** compiler infrastructure as evaluation target
- Can the proposed methodologies:
 - Find new **bugs**? If so, how many?
 - **Improve the performance** of the generated binaries?
 - Improve optimization pass coverage within LLVM?

Bugs

Performance

Coverage

Impacts and Benefits

- **Analysis tools** for compiler developers, to assess their PGO implementations
 - Novel tools to study the problem
 - Enhancement of existing test suites
- **Faster** PGO binaries for the user
 - Reducing the operational costs of deployed applications.
- **Energy savings** due to optimal binaries
 - Reduced environmental footprint.

Thanks for the attention!
Questions?